

Dynamics and Stability of the Spring-Loaded Inverted Pendulum (SLIP) Model

1 Introduction and Relevance

The spring-loaded inverted pendulum (SLIP) is the reduced-order model behind most legged-robot dynamics [3, 5]: a point mass bouncing on a massless spring leg. Despite its simplicity, it captures the center-of-mass dynamics of running in both animals and machines remarkably well [1, 2]. That makes it a good subject for this course: it is simple enough to analyze fully with the tools we covered, but rich enough to produce genuinely interesting hybrid dynamics. It also builds directly on the elastic pendulum studied in class, since the stance phase of SLIP is an inverted spring pendulum.

What makes SLIP interesting from a mechanics standpoint is that it is a hybrid dynamical system: during stance the body behaves like a spring-loaded pendulum, while in flight it is a free projectile, with the two phases connected by touchdown and liftoff switching conditions. Rather than analyzing stability at a fixed equilibrium point, the more natural question for a hopping gait is whether the periodic motion itself is stable. The standard tool for that is the apex return map, a Poincaré map in which a steady hopping gait appears as a fixed point whose stability can be assessed numerically [1, 4].

Beyond being a convenient system, SLIP-type models are used directly in the design and control of running robots and in the biomechanical study of animal and human locomotion, which is my practical motivation for this project: a model simple enough to analyze by hand but predictive enough to inform how a real legged system should be built or controlled.

Relation to Prior Work

This project works within an established line of SLIP research, and each stage of the model draws on a different part of that literature. Blickhan [1] introduced the spring-mass template and showed that it reproduces the ground-reaction-force and center-of-mass patterns of human and animal running; the stance-phase Lagrangian derived here reproduces that same reduced-order model, but from Lagrange's equations in polar coordinates rather than Blickhan's Newtonian force balance. McMahon and Cheng [2] used the same template to relate leg stiffness to running speed empirically; the parameter sweep over stiffness k and touchdown angle θ_{TD} planned for this project asks the same question, but numerically, from the model directly, rather than from experimental data. Full and Koditschek's templates-and-anchors framework [3] is the justification for treating SLIP as a legitimate object of study on its own, rather than only as a shortcut on the way to a fully articulated, multi-link leg model: a template is expected to capture the essential dynamics of the anchored system while discarding detail that does not matter for stability. Schwind and Koditschek [4] derived an approximate, closed-form stance map for SLIP using a perturbative solution; this project instead integrates the stance-phase equations numerically with an event-based solver, trading their analytical insight for an exact (if less interpretable) stance map, consistent with the computational tools emphasized in this course. Saranli, Buehler, and Koditschek [5] built RHex, a hexapod whose individual legs behave like a SLIP template during stance, which is the practical grounding for why this abstraction is worth studying: it is not just a toy system, it is the model actually used to understand a working robot. Finally, Geyer, Seyfarth, and Blickhan [6] extended a SLIP-type compliant-leg model to walking, not just running, by allowing double support; that extension is a natural direction for future work beyond the single-support hopping model considered here.

2 Modeling of Dynamics

The equations of motion for both phases of the SLIP model are derived using **Lagrange's equations of motion**, with generalized coordinates chosen separately for each phase: Cartesian coordinates (x, y) for the flight phase, and polar coordinates (r, θ) about the planted foot for the stance phase. The two phases are connected by touchdown and liftoff switching (transversality) conditions, so the overall model is a hybrid dynamical system integrated with event-based numerical methods. Figure 1 shows the geometry of both phases.




Figure 1: SLIP model geometry: stance (spring pendulum, r, θ) and flight (projectile).

2.1 Flight-Phase Equations of Motion

With basis $\hat{n}_1, \hat{n}_2$, $\vec{r}_m = x\hat{n}_1 + y\hat{n}_2$:

$$T = \frac{1}{2}m(\dot{x}^2 + \dot{y}^2), \quad V = mgy, \quad L = T - V = \frac{1}{2}m(\dot{x}^2 + \dot{y}^2) - mgy$$

Lagrange's equations give

$$\ddot{x} = 0, \quad \ddot{y} = -g$$

Starting at the apex ($\dot{y} = 0$), with $x(0) = x_0$, $\dot{x}(0) = v_{x0}$:

$$x(t) = v_{x0}t + x_0, \quad y(t) = y_0 - \frac{g}{2}t^2$$

2.2 Touchdown: Detection and Transition Map

θ_{TD} is held fixed in flight, so the foot trails the body at

$$\vec{r}_{foot}(t) = (x(t) - r_0 \cos \theta_{TD})\hat{n}_1 + (y(t) - r_0 \sin \theta_{TD})\hat{n}_2$$

Touchdown is $y_{foot} = 0$, i.e. $y_{TD} = r_0 \sin \theta_{TD}$. Solving $y(t) = y_0 - \frac{g}{2}t^2$:

$$t_{TD} = \sqrt{\frac{2(y_0 - r_0 \sin \theta_{TD})}{g}}, \quad \dot{y}_{TD} = -\sqrt{2g(y_0 - r_0 \sin \theta_{TD})}, \quad x_{TD} = v_{x0}t_{TD}, \quad \dot{x}_{TD} = v_{x0}$$

Stance uses polar coordinates about the planted foot, so this Cartesian state must be converted. With foot at $\vec{r}_F = (x_{TD} + r_0 \cos \theta_{TD})\hat{n}_1$ and leg vector $\vec{\rho} = \vec{r}_m - \vec{r}_F = -r_0 \cos \theta_{TD} \hat{n}_1 + y_{TD} \hat{n}_2$:

$$r = |\vec{\rho}| = \sqrt{r_0^2 \cos^2 \theta_{TD} + y_{TD}^2} = r_0, \quad \theta = \text{atan2}(y_{TD}, -r_0 \cos \theta_{TD})$$

Projecting $\vec{v} = \dot{x}_{TD}\hat{n}_1 + \dot{y}_{TD}\hat{n}_2$ onto $\hat{e}_r = \cos\theta\hat{n}_1 + \sin\theta\hat{n}_2$, $\hat{e}_\theta = -\sin\theta\hat{n}_1 + \cos\theta\hat{n}_2$:

$$\dot{r} = \vec{v} \cdot \hat{e}_r = \dot{x}_{TD} \cos\theta + \dot{y}_{TD} \sin\theta, \quad \dot{\theta} = \frac{\vec{v} \cdot \hat{e}_\theta}{r_0} = \frac{-\dot{x}_{TD} \sin\theta + \dot{y}_{TD} \cos\theta}{r_0}$$

implemented as `touchdownState()`.

2.3 Stance-Phase Equations of Motion

In stance the leg pivots about the fixed foot: $\vec{r}_m = r\hat{e}_r$, $\vec{v}_m = \dot{r}\hat{e}_r + r\dot{\theta}\hat{e}_\theta$:

$$T = \frac{1}{2}m\dot{r}^2 + \frac{1}{2}mr^2\dot{\theta}^2, \quad V = mgr \sin\theta + \frac{1}{2}k(r - r_0)^2, \quad L = T - V$$

Lagrange's equations:

$$m\ddot{r} - mr\dot{\theta}^2 + mg \sin\theta + k(r - r_0) = 0, \quad mr^2\ddot{\theta} + 2mr\dot{r}\dot{\theta} + mgr \cos\theta = 0$$

solved for the highest derivatives:

$$\ddot{r} = r\dot{\theta}^2 - g \sin\theta - \frac{k}{m}(r - r_0), \quad \ddot{\theta} = -\frac{2}{r}\dot{r}\dot{\theta} - \frac{g}{r} \cos\theta$$

implemented as `standingState()`.

2.4 Conserved Energy in Stance

$\vec{r}_m = r\hat{e}_r$ is scleronomic and T is homogeneous quadratic in $(\dot{r}, \dot{\theta})$, so $H = T + V$; with no explicit t -dependence, energy is conserved through stance:

$$E = \frac{1}{2}m(\dot{r}^2 + r^2\dot{\theta}^2) + mgr \sin\theta + \frac{1}{2}k(r - r_0)^2 = \text{const.}$$

Neither coordinate is cyclic, so both equations must be integrated together; this conserved quantity is used below as a check on the numerical integrator.

2.5 Liftoff Map and the Apex Return Map

Liftoff occurs at $r = r_0$, $\dot{r} > 0$. Chaining flight $\rightarrow$ touchdown $\rightarrow$ stance $\rightarrow$ liftoff $\rightarrow$ flight once defines the apex-to-apex map

$$P : (y_n, v_{x,n}) \mapsto (y_{n+1}, v_{x,n+1}),$$

whose fixed points are steady periodic hopping gaits and whose stability is governed by the linearization of P there. Section 2.7 reduces this to a single scalar, so that linearization is one derivative rather than a Jacobian.

Liftoff position converts to Cartesian from $\vec{r}_m = r\hat{e}_r$ at $r = r_0$, $\theta = \theta_{LO}$:

$$x_{LO} = r_0 \cos\theta_{LO}, \quad y_{LO} = r_0 \sin\theta_{LO}$$

Velocity converts by substituting $\hat{e}_r = \cos\theta\hat{n}_1 + \sin\theta\hat{n}_2$, $\hat{e}_\theta = -\sin\theta\hat{n}_1 + \cos\theta\hat{n}_2$ into $\vec{v}_m = \dot{r}\hat{e}_r + r\dot{\theta}\hat{e}_\theta$ and evaluating at liftoff:

$$\dot{x}_{LO} = \dot{r}_{LO} \cos \theta_{LO} - r_0 \dot{\theta}_{LO} \sin \theta_{LO}, \quad \dot{y}_{LO} = \dot{r}_{LO} \sin \theta_{LO} + r_0 \dot{\theta}_{LO} \cos \theta_{LO}$$

Once airborne, apex is defined in the ground frame by $\dot{y} = 0$. With $\ddot{y} = -g$ and the clock started at liftoff, $\dot{y}(t) = \dot{y}_{LO} - gt$, so

$$t_{\text{apex}} = \frac{\dot{y}_{LO}}{g}$$

Propagating the flight solution $x(t) = \dot{x}_{LO}t + x_{LO}$, $y(t) = y_{LO} + \dot{y}_{LO}t - \frac{g}{2}t^2$ to $t = t_{\text{apex}}$:

$$x_{n+1} = x_{LO} + \dot{x}_{LO} t_{\text{apex}}, \quad y_{n+1} = y_{LO} + \dot{y}_{LO} t_{\text{apex}} - \frac{g}{2} t_{\text{apex}}^2, \quad v_{x,n+1} = \dot{x}_{LO}, \quad \dot{y}_{n+1} = 0$$

Both are exact: $v_{x,n+1} = \dot{x}_{LO}$ since horizontal motion is unforced in flight ($\ddot{x} = 0$), and $\dot{y}_{n+1} = 0$ by construction of t_{apex} . x_{LO} is measured relative to the stance foot just released rather than a fixed ground origin. This is consistent, since x never enters the dynamics of either phase (Section 2.7).

2.6 Conserved Energy at Apex and Across a Full Hop Cycle

By the same argument as Section 2.4, the flight Lagrangian is scleronomic with T homogeneous quadratic in $(\dot{x}, \dot{y})$, so $H = T + V$ is conserved through flight:

$$E = \frac{1}{2}m(\dot{x}^2 + \dot{y}^2) + mgy = \text{const.}$$

with no spring term, since the leg carries no elastic energy while airborne. This differs from the stance energy $E_{\text{stance}} = \frac{1}{2}m(\dot{r}^2 + r^2\dot{\theta}^2) + mgr \sin \theta + \frac{1}{2}k(r - r_0)^2$ only by that spring term, which vanishes at $r = r_0$ – exactly the touchdown and liftoff condition. The two functionals therefore agree at every switching event, so total mechanical energy is continuous across touchdown and liftoff, and conserved over the full cycle apex $\rightarrow$ touchdown $\rightarrow$ stance $\rightarrow$ liftoff $\rightarrow$ apex. Section 4.2 validates this numerically.

2.7 Dimension Reduction of the Apex State

At apex, $\dot{y} = 0$ by definition. Substituting into the apex energy of Section 2.6 and solving for $\dot{x}$:

$$E - mgy = \frac{1}{2}m\dot{x}^2 \quad \implies \quad \dot{x}^2 = \frac{2(E - mgy)}{m} = \frac{2E}{m} - 2gy$$

$$\dot{x} = \sqrt{\frac{2E}{m} - 2gy} \quad (\dot{x} > 0 \text{ for forward hopping})$$

So the apex state $(x, y, \dot{x}, \dot{y})$ is not four independent numbers: $\dot{y} = 0$ always, x is ignorable (it enters no potential or force in either phase), and $\dot{x}$ is fixed by y once E is fixed. At fixed E , the apex state collapses to the scalar y , and the hop-to-hop dynamics reduce to the one-dimensional map

$$y_{n+1} = P(y_n)$$

where P is the composition: reconstruct $\dot{x}_n$ from y_n (above) $\rightarrow$ flight to touchdown (Section 2.2) $\rightarrow$ stance to liftoff (Section 2.3) $\rightarrow$ flight to the next apex (Section 2.5). Section 4.2 confirms this reduction numerically: $\dot{x}$ evaluated at the simulated output height y_{n+1} matches the simulator's own $\dot{x}_{n+1}$ to machine precision.

2.8 The Apex Return Map: Fixed Point and Stability Criterion

Using the reduction of Section 2.7, define the apex return map P at fixed leg stiffness k , fixed touchdown angle θ_{TD} , and fixed total energy E :

$$P : y_n \mapsto y_{n+1}$$

E characterizes how energetic a gait is; k and θ_{TD} are the model's only free parameters. There is no controller, so a periodic gait at a given E is whatever P produces for a chosen (k, θ_{TD}) .

A periodic gait is a fixed point $y^* = P(y^*)$, equivalently a root of

$$f(y) = P(y) - y$$

Local stability is governed by the derivative at the fixed point, estimated by a central finite difference:

$$P'(y^*) \approx \frac{P(y^* + dy) - P(y^* - dy)}{2 dy}$$

$$|P'(y^*)| < 1 : \text{stable}, \quad |P'(y^*)| > 1 : \text{unstable}$$

$P'(y^*) < 0$ indicates oscillatory divergence/convergence (perturbations alternate sides of y^*); $P'(y^*) > 0$ indicates monotonic divergence/convergence. Section 4.3 reports the root-finding methodology and the result for the validated parameter set of Section 4.1.

3 Discussion of Assumptions

Several modeling choices were made to keep the SLIP model tractable with the tools of this course, each with a corresponding limitation:

- **Planar motion.** The model is restricted to the sagittal plane. Real running involves some lateral and rotational motion of the trunk and legs, which this model cannot capture.
- **Point mass, massless leg.** All body mass is lumped at a single point, and the leg itself carries no mass or rotational inertia. This isolates the center-of-mass dynamics that SLIP is meant to describe, but discards any effect of leg-swing inertia or trunk pitch, which the templates-and-anchors view [3] explicitly treats as detail that a template is allowed to omit.
- **Linear (Hookean) leg spring.** The stance leg is modeled as a single ideal linear spring, $F = -k(r - r_0)$, with constant stiffness k . Biological and robotic legs have leg stiffness that is nonlinear and can be actively modulated by the neuromuscular system or a controller; treating k as a fixed parameter to be swept is a simplification.
- **Fixed touchdown angle held through flight.** θ_{TD} is assumed constant for the entire flight phase, i.e. there is no swing-leg retraction or angle adjustment in response to a disturbance before the foot lands. This is the standard SLIP assumption, but it means the model's only "control input" between hops is the touchdown angle chosen before each flight, not a continuous correction during flight.
- **Instantaneous, rigid, non-slipping foot placement.** Touchdown and liftoff are treated as instantaneous events with no impact losses, and the foot is assumed to stick to the ground (no slipping, no ground compliance) for the duration of stance at a fixed height $y = 0$. Real contact involves finite compliance and a friction cone that could, in principle, be violated at high touchdown angles or high speeds.

- **No stance-phase dissipation, no aerodynamic drag.** The stance-phase Lagrangian has no damping term, so mechanical energy is exactly conserved within a single stance phase (verified numerically in Section 4); flight is likewise treated as drag-free projectile motion. This was a deliberate choice: it isolates the elastic dynamics of interest and gives a clean, near-machine-precision check on the numerical integrator, at the cost of neglecting the fact that any real leg or running surface dissipates some energy each stride.
- **Equal rest length at touchdown and liftoff.** The leg is assumed to return to the same natural length r_0 between hops, with no active leg extension/retraction actuation changing that rest length from one stance phase to the next.

4 Computational Results

4.1 Single-Stance Validation

The model is implemented in Python. `main.py` maps an apex state to the stance-phase initial condition via `touchdownState()`, then integrates `standingState()` with `solve_ivp` (`rtol` = 10^{-10} , `atol` = 10^{-12}) to the liftoff event $r = r_0$, $\dot{r} > 0$, checking the conserved energy from Section 2.4 at every step. Using the parameters below, this validates the stance-phase model end to end:

m	r_0	k	g	θ_{TD}	y_0	v_{x0}
80 kg	1 m	20000 N/m	9.81 m/s ²	68°	1.2 m	3.0 m/s

	r [m]	θ [deg]	$\dot{r}$ [m/s]	$\dot{\theta}$ [rad/s]
Touchdown	1.0000	112.00	−3.2689	−1.9149
Liftoff	1.0000	83.48	3.1937	−1.6886

Stance lasts 0.2094 s; the leg compresses to 0.7740 m (22.6% of r_0); relative energy drift is 6.46×10^{-13} , i.e. machine precision, confirming that the integrator and the equations of motion in Section 2.3 are consistent. Figure 2 plots leg length, angle, CoM trajectory, and energy drift for this run.




Figure 2: Stance simulation: $r(t)$, $\theta(t)$ (top), CoM trajectory (bottom left), energy drift (bottom right).

4.2 Full-Cycle Energy Conservation and Dimension-Reduction Validation

`proveX_dotIsRedundant()` validates the full hop cycle: starting from the apex state of Section 4.1, it computes the apex energy E_1 (`energyHop()`), chains one hop – flight $\rightarrow$ touchdown $\rightarrow$ stance $\rightarrow$ liftoff $\rightarrow$ flight (`hop()`) – and recomputes the apex energy E_2 . Relative drift over the full cycle,

$$\frac{|E_2 - E_1|}{|E_1|} = 6.17 \times 10^{-13},$$

is consistent with the stance-only drift of Section 4.1 (6.46×10^{-13}), confirming total mechanical energy is conserved across the full switching cycle, not just within stance.

The same run checks the dimension reduction of Section 2.7: with E_1 fixed, $\dot{x} = \sqrt{2E_1/m - 2gy}$ evaluated at the simulated output height $y_{n+1} = 1.4465$ m predicts $\dot{x}_{pred} = 2.0403$ m/s, matching the simulator's own $v_{x,n+1} = 2.0403$ m/s to

$$\frac{|v_{x,n+1} - \dot{x}_{pred}|}{|\dot{x}_{pred}|} = 2.41 \times 10^{-12}.$$

4.3 Apex Return Map: Root-Finding Methodology and Result

`findApexMapFixedPoint()` locates the root of $f(y) = P(y) - y$ (Section 2.8) over the domain of y physically valid for a given (k, θ_{TD}, E) : bounded below by the minimum apex height consistent with touchdown at angle θ_{TD} ,

$$y_{min} = r_0 \sin \theta_{TD},$$

and above by the height at which the closed-form $\dot{x}$ of Section 2.7 would require a negative value under the square root,

$$y_{max} = \frac{E}{mg}.$$

A small inward padding ($\epsilon = 10^{-3}$ m) is added to y_{min} to avoid floating-point roundoff pushing the touchdown free-fall relation's square root slightly negative at the exact boundary.

$f(y)$ is sampled coarsely (30 points) across $[y_{min}, y_{max}]$; once a bracket with a sign change is found, Brent's method refines it to a root y^* , and $P'(y^*)$ is estimated by the central difference of Section 2.8 with $dy = 10^{-3}$ m.

For the validated parameter set of Section 4.1 ($m = 80$ kg, $r_0 = 1$ m, $k = 20000$ N/m, $\theta_{TD} = 68^\circ$, E fixed by $y_0 = 1.2$ m, $v_{x0} = 3.0$ m/s):

$$y^* \approx 1.5096 \text{ m}, \quad P'(y^*) \approx -2.76$$

$|P'(y^*)| > 1$: this fixed point is unstable, and the negative sign indicates oscillatory divergence – a perturbed apex height overshoots on alternating sides of y^* , growing each hop. Steady hopping at $\theta_{TD} = 68^\circ$, $k = 20000$ N/m is therefore not stable.

4.4 Grid Sweep Methodology and Parameter Sweep

`gridSearchFixedPoints()` wraps the pipeline of Section 4.3 into a function of (k, θ_{TD}) returning $(y^*, P'(y^*))$, or null if no sign change in $f(y)$ is found – i.e. no periodic gait exists there at this E . Each cell uses its own copy of the parameter dictionary, so sweeping k never mutates state shared across cells. `stableBandWidth()` additionally reports the longest run of contiguous stable cells along each grid direction, which is used below to tell a genuinely narrow stability margin from a coarse-grid artifact.

An initial coarse 40×40 diagnostic grid over $k \in [5000, 40000]$ N/m, $\theta_{TD} \in [55^\circ, 80^\circ]$ turned up only a one-cell-wide staircase of stable cells (band width 2 cells along θ_{TD} , 3 along k) threading through an otherwise unstable field – too thin to trust at that resolution. `localRefinementSweep()` zoomed a 20×20 window ($\pm 2.5^\circ \times \pm 2500$ N/m, i.e. roughly $4 \times$ finer) onto a representative stable cell near $k \approx 20256$ N/m, $\theta_{TD} \approx 71.0^\circ$; there the stable band *thickened* (3 cells along θ_{TD} , 5 along k), meaning the coarse staircase was a grid-resolution artifact rather than the true boundary shape, so the full sweep needed that finer spacing before its shape could be trusted.

The final grid matches that resolution across the whole domain: a 133×95 sweep ($k \in [5000, 40000]$ N/m in steps of ≈ 265 N/m, $\theta_{TD} \in [55^\circ, 80^\circ]$ in steps of $\approx 0.27^\circ$; 12,635 cells, of which only 5 found no fixed point), classified stable/unstable/no-fixed-point in Figure 3, with the friction-cone floor from Section 4.3 ($\theta_{TD} \geq 59^\circ$, $\mu = 0.6$) overlaid. At this resolution the band widens slightly (3 cells along θ_{TD} , 8 along k) but is still only a few cells across. A second local refinement at the same fine spacing (Figure 4), now centered on the final grid's own representative stable cell ($k \approx 20114$ N/m, $\theta_{TD} \approx 70.96^\circ$, where $y^* \approx 0.9496$ m and $P'(y^*) \approx 0.348$), gives a band width of 3 cells along θ_{TD} and 5 along k – the same order as the full grid, not larger, so the earlier thickening has stopped: the self-stable region is a real, if narrow, sliver of (k, θ_{TD}) space at this E , not an artifact of grid spacing.




Figure 3: Self-stable running region over the full 133×95 (k, θ_{TD}) grid at fixed E (Section 4.1 parameters). Green: stable fixed point ($|P'(y^*)| < 1$); orange: unstable; gray: no fixed point found. Dashed line/shading: the friction-cone floor $\theta_{TD} \geq 59^\circ$ ($\mu = 0.6$).




Figure 4: Local refinement ($20 \times 20, \pm 2.5^\circ \times \pm 2500$ N/m) around the stable band near $k \approx 20114$ N/m, $\theta_{TD} \approx 70.96^\circ$, at the same fine spacing used to validate Figure 3's resolution. The band's width here matches the full grid, confirming it is not still thickening.

The sweep and its resolution check are done; what remains is the write-up:

- An iteration plot (e.g. y_{n+1} vs. y_n , or y_n vs. hop index n) showing convergence to a fixed point for a *stable* case – the representative cell above ($k \approx 20114$ N/m, $\theta_{TD} \approx 70.96^\circ$, $y^* \approx 0.9496$ m, $P'(y^*) \approx 0.348$) is a ready-made example, and since $P'(y^*) > 0$ it should converge *monotonically*, in contrast to the oscillatory divergence of Section 4.3’s unstable case.
- Interpretation: does the shape of the stable band (stiffness rising with touchdown angle) match the qualitative stiffness–speed trend reported by McMahon and Cheng [2]? Section 4.3’s point (68° , 20000 N/m) is unstable and sits just outside the band at that stiffness – is that representative of the grid, or an outlier chosen before the sweep existed? What does it mean physically that the self-stable region is only a few cells wide even at full resolution – is steady hopping in this model fundamentally fragile, or is this an artifact of holding θ_{TD} fixed through flight (Section 3) with no swing-leg feedback?

5 Lessons Learned

This project proved to be incredibly informative. It allowed me to dive into a multi-stage dynamic system and taught me how to address the state transition conditions in a clear and defined way. It also expanded upon our study of geometric dynamics analysis by requiring a different mechanism for the stability comparison. Instead of having to analyze the whole system, we only have to look at the points in space that best define our motion. This also drove home that finding a periodic solution is not the same as finding a stable one: the fixed point of the apex return map exists whether or not it is attracting, so stability has to be checked separately – and for the validated parameters, it turned out the hopping gait was not stable at all.

It was also a great project for working through robotics-style problems. The inverted pendulum is a very famous model used throughout all of robotics. But with the Lagrangian, this derivation can become relatively simple, opening the door for more accurate models. Working through the parameter sweep taught me a practical lesson about trusting numerical results, too: the first coarse grid over (k, θ_{TD}) turned up a stable band so thin it could have just been a grid-resolution artifact, and only after refining the grid and confirming the band held its width could I be confident that the narrow self-stable region was a real feature of the model rather than a numerical accident.

References

- [1] R. Blickhan, “The spring-mass model for running and hopping,” *Journal of Biomechanics*, vol. 22, no. 11–12, pp. 1217–1227, 1989.
- [2] T. A. McMahon and G. C. Cheng, “The mechanics of running: how does stiffness couple with speed?,” *Journal of Biomechanics*, vol. 23, suppl. 1, pp. 65–78, 1990.
- [3] R. J. Full and D. E. Koditschek, “Templates and anchors: neuromechanical hypotheses of legged locomotion on land,” *Journal of Experimental Biology*, vol. 202, no. 23, pp. 3325–3332, 1999.
- [4] W. J. Schwind and D. E. Koditschek, “Approximating the stance map of a 2-DOF monopod runner,” *Journal of Nonlinear Science*, vol. 10, pp. 533–568, 2000.
- [5] U. Saranli, M. Buehler, and D. E. Koditschek, “RHex: a simple and highly mobile hexapod robot,” *International Journal of Robotics Research*, vol. 20, no. 7, pp. 616–631, 2001.
- [6] H. Geyer, A. Seyfarth, and R. Blickhan, “Compliant leg behaviour explains basic dynamics of walking and running,” *Proceedings of the Royal Society B*, vol. 273, no. 1603, pp. 2861–2867, 2006.

Appendix: Code

main.py – touchdown map, stance-phase RHS, and the single-stance simulation and validation plots discussed in Section 4.1.

```

1 import copy
2 import time
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from scipy.integrate import solve_ivp
6 from scipy.optimize import least_squares
7 from scipy.optimize import brentq
8
9 # Paramerts
10 m = 80
11 r_o = 1
12 k = 20000
13 g = 9.81
14 thetaTD = np.radians(68)
15 y_o = 1.2
16 v_o = 3.0
17
18 params = {'m' : m , 'k' : k , 'g' : g , 'r_o' : r_o , 'y_o' : y_o}
19
20
21 # =====
22 # Stage 1 - Stance phase: touchdown -> standing -> liftoff
23 # First piece of the model: integrate a single stance phase and
24 # check that total mechanical energy is conserved through it.
25 # =====
26
27 def touchdownState(x_o_dot , y_o , theta_TD , r_o , g):
28     r = r_o
29     theta = np.arctan2(r_o * np.sin(theta_TD) , -r_o * np.cos(theta_TD))
30
31     x_TD_dot = x_o_dot
32     y_TD_dot = - np.sqrt(2*g*(y_o - r_o * np.sin(theta_TD)))
33
34     r_dot = x_TD_dot * np.cos(theta) + y_TD_dot * np.sin(theta)
35

```

```

36     theta_dot = (-x_TD_dot * np.sin(theta) + y_TD_dot * np.cos(theta)) / r_o
37
38     return [r , theta, r_dot , theta_dot]
39
40 def standingState(t , s, params):
41     r_dot = s[2]
42     theta_dot = s[3]
43
44     r_ddot = s[0] * theta_dot**2 - params['g'] * np.sin(s[1]) - (params['k'] /
45         params['m'])* (s[0] - params['r_o'])
46     theta_ddot = -(2/s[0]) * r_dot * theta_dot - (params['g'] / s[0]) * np.cos
47         (s[1])
48
49     return [r_dot , theta_dot, r_ddot, theta_ddot]
50
51 def liftoffStateSwitch(t, s, params):
52     return params['k'] * (s[0] - params['r_o'])
53 liftoffStateSwitch.terminal = True
54 liftoffStateSwitch.direction = 1
55
56 def energy(s, params):
57     E = (1/2) * params['m'] * (s[2]**2 + s[0]**2 * s[3]**2) + params['m'] *
58         params['g'] * s[0] * np.sin(s[1]) + (1/2) * params['k'] * (s[0] -
59         params['r_o'])**2
60
61     return E
62
63 # ----- run a single stance phase for diagnostics -----
64 # (mirrors the body of hop(), but keeps sol/s0 around for plotting)
65
66 def initialFunctionTest():
67     s0 = touchdownState(v_o, y_o, thetaTD, r_o, g)
68
69     t_span = (0.0, 1.0)
70     sol = solve_ivp(standingState, t_span, s0, args=(params,),
71         events=liftoffStateSwitch, rtol=1e-10, atol=1e-12,
72         dense_output=True, max_step=1e-3)
73
74     E = np.array([energy(sol.y[:, i], params) for i in range(sol.y.shape[1])])
75     drift = np.max(np.abs(E - E[0])) / np.abs(E[0])
76
77     # Plotting and data stuff, AI can fill in these gaps as simply learning to
78     # plot good is not hard and not my current goal
79
80     print("status:", sol.status)
81     print("liftoffStateSwitch time:", sol.t_events[0])
82     print("drift:", drift)
83
84     # ----- diagnostics -----
85     r_TD, th_TD, rdot_TD, thdot_TD = s0
86     s_LO = sol.y_events[0][0]
87     t_LO = sol.t_events[0][0]
88
89     print("--- touchdown state ---")
90     print(f"    r          = {r_TD:.6f} m")
91     print(f"    theta      = {np.degrees(th_TD):.4f} deg")

```

```

87     print(f"   r_dot   = {rdot_TD:.6f} m/s")
88     print(f"   th_dot  = {thdot_TD:.6f} rad/s")
89
90     print("--- liftoffStateSwitch state ---")
91     print(f"   status      = {sol.status}")
92     print(f"   stance duration= {t_L0:.6f} s")
93     print(f"   r            = {s_L0[0]:.6f} m")
94     print(f"   theta        = {np.degrees(s_L0[1]):.4f} deg")
95     print(f"   r_dot        = {s_L0[2]:.6f} m/s")
96     print(f"   th_dot       = {s_L0[3]:.6f} rad/s")
97
98     print("--- validation ---")
99     print(f"   min leg length = {np.min(sol.y[0]):.6f} m")
100    print(f"   max compression= {params['r_o'] - np.min(sol.y[0]):.6f} m")
101    print(f"   E_0           = {E[0]:.6f} J")
102    print(f"   relative drift = {drift:.3e}")
103
104    # ----- figures -----
105    t = sol.t
106    r = sol.y[0]
107    th = sol.y[1]
108    x_com = r * np.cos(th)
109    y_com = r * np.sin(th)
110
111    fig, ax = plt.subplots(2, 2, figsize=(11, 8))
112
113    ax[0, 0].plot(t, r, color='C3')
114    ax[0, 0].axhline(params['r_o'], ls='--', lw=0.8, color='gray', label=
115        '$r_0$')
116    ax[0, 0].set_xlabel('t [s]')
117    ax[0, 0].set_ylabel('r [m]')
118    ax[0, 0].set_title('Leg length')
119    ax[0, 0].legend()
120
121    ax[0, 1].plot(t, np.degrees(th), color='C0')
122    ax[0, 1].axhline(90, ls='--', lw=0.8, color='gray', label='$90^\circ$')
123    ax[0, 1].set_xlabel('t [s]')
124    ax[0, 1].set_ylabel(r'$\theta$ [deg]')
125    ax[0, 1].set_title('Leg angle')
126    ax[0, 1].legend()
127
128    ax[1, 0].plot(x_com, y_com, color='C3', lw=2)
129    ax[1, 0].plot(0, 0, 'kv', ms=10, label='foot (pinned)')
130    ax[1, 0].plot(x_com[0], y_com[0], 'o', color='C0', label='touchdown')
131    ax[1, 0].plot(x_com[-1], y_com[-1], 's', color='C2', label='
132        liftoffStateSwitch')
133    ax[1, 0].axhline(0, color='k', lw=1)
134    ax[1, 0].set_xlabel('x [m]')
135    ax[1, 0].set_ylabel('y [m]')
136    ax[1, 0].set_title('CoM trajectory (foot frame)')
137    ax[1, 0].axis('equal')
138    ax[1, 0].legend()
139
140    ax[1, 1].plot(t, E - E[0], color='C4')
141    ax[1, 1].set_xlabel('t [s]')
142    ax[1, 1].set_ylabel('$E - E_0$ [J]')

```

```

141 ax[1, 1].set_title(f'Energy drift (rel. {drift:.2e})')
142 ax[1, 1].ticklabel_format(axis='y', style='sci', scilimits=(0, 0))
143
144 for a in ax.flat:
145     a.grid(alpha=0.3)
146 fig.tight_layout()
147 fig.savefig(os.path.join(Script_DIR, 'stance_validation.png'), dpi=200)
148 plt.show()
149
150
151 # =====
152 # Stage 2 - Apex state: ballistic flight after liftoff
153 # Convert the liftoff (r, theta) state back to Cartesian (x, y)
154 # and propagate the free-flight trajectory forward to the apex
155 # (the point where the vertical velocity goes to zero).
156 # =====
157
158 def projectileApexState(s):
159     r_o = s[0]
160     theta_o = s[1]
161     r_o_dot = s[2]
162     theta_o_dot = s[3]
163
164     # convert state from liftoff back to cartesian
165     y_o = r_o * np.sin(theta_o)
166     x_o = r_o * np.cos(theta_o)
167
168     y_o_dot = r_o_dot * np.sin(theta_o) + theta_o_dot * r_o * np.cos(theta_o)
169     x_o_dot = r_o_dot * np.cos(theta_o) - theta_o_dot * r_o * np.sin(theta_o)
170
171     # Solve for when y velocity = 0, this is the highest point
172     t_apex = y_o_dot / g
173
174     # find corresponding x y x_dot y_dot values
175     y_dot = -g*t_apex + y_o_dot
176     x_dot = x_o_dot
177
178     y = y_o - ((g/2) * t_apex**2) + y_o_dot * t_apex
179     x = x_o + x_dot * t_apex
180
181     return [x, y, x_dot, y_dot]
182
183
184 # =====
185 # Stage 3 - Hop: touchdown -> stance -> liftoff -> apex
186 # Chains stages 1-2 together into a single map that takes an
187 # apex state and returns the next apex state, one hop later.
188 # =====
189
190 def hop(apexState, thetaTD, params):
191     s0 = touchdownState(apexState[2], apexState[1], thetaTD, params['r_o'],
192                          params['g'])
193     t_span = (0.0, 1.0)
194
195     sol = solve_ivp(standingState, t_span, s0, args=(params,),
196                     events=liftoffStateSwitch, rtol=1e-10, atol=1e-12,

```

```

196         dense_output=True, max_step=1e-3)
197
198     s1 = projectileApexState(sol.y_events[0][0])
199
200     return s1
201
202 def energyHop(apexState, params):
203     E = (1/2) * params['m'] * (apexState[2]**2 + apexState[3]**2) + params['m',
204         ] * params['g'] * apexState[1]
205     return E
206
207 # =====
208 # Stage 4 - x is redundant: the apex-to-apex map is fully
209 # determined by total energy E and height y, not by x directly.
210 # =====
211
212 def proveX_dotIsRedundant():
213     # E = (1/2)*m*(x_dot^2 + y_dot^2) + m*g*y
214     # ((E - mgy) * 2/m - y_dot**2)**.5 = x_dot
215     # Y_dot = 0
216     # ((E-mgy)*2/m)**.5 = x_dot
217     # (2E/m - 2gy)**.5 = x_dot
218
219     initialState = [0,y_o,v_o,0]
220     initialThetaTD = thetaTD
221
222     E1 = energyHop(initialState, params)
223
224
225     finalState = hop(initialState, initialThetaTD, params)
226     x_dot_pred = np.sqrt(2*E1/ params['m'] - 2*params['g']*finalState[1])
227     E2 = energyHop(finalState, params)
228
229     print(f"Energy Drift over 1 hop: {E2-E1}")
230     print(f"X_dot Drift over 1 hop: {finalState[2]-x_dot_pred}")
231
232 # =====
233 # Stage 5 - Reduced 1-D apex map: fixed point and stability
234 # Stage 4 showed the hop only depends on E and y, so collapse the
235 # hop into a 1-D map y -> apexmap(y) and find the periodic apex
236 # height (fixed point), then check its local stability from the
237 # map's slope there.
238 # =====
239
240 def apexmap(y, thetaTD, params, E):
241     y = np.asarray(y).item()
242
243     x_dot = np.sqrt(2*E/ params['m'] - 2*params['g']*y)
244
245     inputApexState = [0, y, x_dot, 0]
246     outputApexState = hop(inputApexState, thetaTD, params)
247
248     return outputApexState[1]
249
250 # ----- find the fixed point of the apex map and its local stability

```

```

-----
251
252 def findApexMapFixedPoint(params, thetaTD, y_o=y_o, v_o=v_o):
253     initialState = [0,y_o,v_o,0]
254     E = energyHop(initialState, params)
255     f = lambda y: apexmap(y, thetaTD, params, E) - y
256
257     # +.001 avoids singularities
258     y_min = params['r_o'] * np.sin(thetaTD) + .001 # apex must be at/above
           touchdown height for this thetaTD
259     #  $2E/m - 2*g*y \geq 0$ 
260     #  $y < 2E / (2mg)$ 
261     y_max = 2*E / (2* params['m'] * params['g'])
262
263     y_sample = np.linspace(y_min, y_max, 30)
264     f_sample = np.array([f(yi) for yi in y_sample])
265
266     sign_changes = np.where(np.diff(np.sign(f_sample)) != 0)[0]
267
268     if sign_changes.size == 0:
269         raise ValueError(f"no sign change in apex map for thetaTD={thetaTD}, k
           ={{params['k']}}")
270
271     y_min = y_sample[sign_changes[0]]
272     y_max = y_sample[sign_changes[0] + 1]
273
274     y_star: float = brentq(f, y_min, y_max) # type: ignore[assignment]
275
276     dy = max(1e-4 * abs(y_star), 1e-6)
277
278     apexmapPrime = (apexmap(y_star + dy, thetaTD, params, E) - apexmap(y_star
           - dy, thetaTD, params, E)) / (2 * dy) # type: ignore
279
280     # print(apexmapPrime)
281
282     return (y_star, apexmapPrime)
283
284
285 # =====
286 # Stage 6 - Grid search: fixed point and stability over (k, thetaTD)
287 # Each grid point needs its own params dict (deep-copied off the
288 # base params) so sweeping 'k' never mutates the shared dict that
289 # other stages/functions read from.
290 # =====
291
292 def gridSearchFixedPoints(base_params, k_values, thetaTD_values, y_o=y_o, v_o=
v_o):
293     y_star_grid = np.full((len(k_values), len(thetaTD_values)), np.nan)
294     slope_grid = np.full((len(k_values), len(thetaTD_values)), np.nan)
295     n_failed = 0
296     n_total = len(k_values) * len(thetaTD_values)
297
298     # High-resolution sweeps can take well over an hour unattended, so print
299     # a per-row progress/ETA line rather than going silent until the end.
300     t_start = time.time()
301

```



```

302     for i, k_val in enumerate(k_values):
303         for j, theta_val in enumerate(thetaTD_values):
304             trial_params = copy.deepcopy(base_params)
305             trial_params['k'] = k_val
306             try:
307                 y_star, slope = findApexMapFixedPoint(trial_params, theta_val,
308                                                         y_o, v_o)
309             except ValueError:
310                 n_failed += 1
311                 continue
312             y_star_grid[i, j] = y_star
313             slope_grid[i, j] = slope
314
315         n_done = (i + 1) * len(thetaTD_values)
316         elapsed = time.time() - t_start
317         eta = elapsed * (n_total - n_done) / n_done
318         print(f"    row {i+1}/{len(k_values)} (k={k_val:.0f}) done -- "
319               f"{n_done}/{n_total} cells, {elapsed/60:.1f} min elapsed, "
320               f"~{eta/60:.1f} min remaining")
321
322     print(f"{n_failed}/{n_total} cells failed to find a fixed point")
323     return y_star_grid, slope_grid
324
325
326 import numpy as np
327 import matplotlib.pyplot as plt
328 from matplotlib.colors import ListedColormap, BoundaryNorm
329
330 def plotStabilitySweep(k_values, thetaTD_values, y_star_grid, slope_grid,
331                       friction_floor_deg=59.0, save_path=None):
332     """
333     Classifies each grid cell as stable / unstable / no fixed point,
334     overlays the friction-cone floor, and saves the deliverable figure.
335     """
336     K, THETA = np.meshgrid(k_values, thetaTD_values, indexing='ij')
337
338     # 0 = no fixed point, 1 = unstable, 2 = stable
339     classification = np.zeros_like(slope_grid)
340     found = ~np.isnan(slope_grid)
341     classification[found & (np.abs(slope_grid) < 1)] = 2
342     classification[found & (np.abs(slope_grid) >= 1)] = 1
343
344     cmap = ListedColormap(['#d9d9d9', '#e07b54', '#4c9a6a']) # gray, orange,
345     green
346     bounds = [-0.5, 0.5, 1.5, 2.5]
347     norm = BoundaryNorm(bounds, cmap.N)
348
349     fig, ax = plt.subplots(figsize=(8, 6))
350     mesh = ax.pcolormesh(THETA, K, classification, cmap=cmap, norm=norm,
351                          shading='auto')
352
353     # friction-cone floor
354     ax.axvline(friction_floor_deg, color='k', linestyle='--', linewidth=1.5)
355     ax.axvspan(thetaTD_values.min(), friction_floor_deg, color='k', alpha
356               =0.08)

```

```

354 ax.text(friction_floor_deg + 0.5, k_values.max() * 0.97,
355         r'$\theta_{TD} \geq %.0f^\circ$ (friction cone, $\mu=0.6$)' %
356         friction_floor_deg,
357         va='top', fontsize=9)
358
359 ax.set_xlabel(r'Touchdown angle $\theta_{TD}$ (deg)')
360 ax.set_ylabel(r'Leg stiffness $k$')
361 ax.set_title('Self-stable running region')
362
363 cbar = fig.colorbar(mesh, ax=ax, ticks=[0, 1, 2])
364 cbar.ax.set_yticklabels(['No fixed point', 'Unstable', 'Stable'])
365
366 fig.tight_layout()
367 if save_path:
368     fig.savefig(save_path, dpi=300)
369 return fig, ax
370
371 # =====
372 # Stage 7 - Local refinement:
373 # =====
374
375 def localRefinementSweep(base_params, k_center, theta_center_deg,
376                          k_half_width=2500.0, theta_half_width_deg=2.5,
377                          n=20, y_o=y_o, v_o=v_o):
378     """Rerun a small (k, thetaTD) window at finer resolution, zoomed on a
379     stable cell found by the coarse sweep."""
380     k_values = np.linspace(k_center - k_half_width, k_center + k_half_width, n)
381     thetaTD_deg = np.linspace(theta_center_deg - theta_half_width_deg,
382                               theta_center_deg + theta_half_width_deg, n)
383     thetaTD_values = np.radians(thetaTD_deg)
384
385     y_star_grid, slope_grid = gridSearchFixedPoints(base_params, k_values,
386                                                     thetaTD_values, y_o, v_o)
387     return k_values, thetaTD_deg, y_star_grid, slope_grid
388
389 def stableBandWidth(slope_grid):
390     """Longest run of contiguous stable ( $|\text{slope}| < 1$ ) cells found along any
391     single row (fixed k, varying thetaTD) or column (fixed thetaTD, varying
392     k). A width of 1 on a fine grid means the band really is that thin at
393     this resolution; >1 means the coarse sweep's one-cell staircase was an
394     artifact of the coarse spacing."""
395     stable = (~np.isnan(slope_grid)) & (np.abs(slope_grid) < 1)
396
397     def longest_run(mask_1d):
398         best = run = 0
399         for is_stable in mask_1d:
400             run = run + 1 if is_stable else 0
401             best = max(best, run)
402         return best
403
404     row_runs = [longest_run(row) for row in stable]
405     col_runs = [longest_run(col) for col in stable.T]
406     return max(row_runs, default=0), max(col_runs, default=0)

```

```

407
408
409 # =====
410 # Run
411 # =====
412
413 import os
414 import warnings
415
416 SCRIPT_DIR = os.path.dirname(os.path.abspath(__file__))
417
418 RESULTS_PATH = os.path.join(SCRIPT_DIR, "sweep_results.npz")
419 RERUN_SWEEP = False # set True to force a fresh sweep even if a cached file
    exists
420
421
422 def runOrLoadSweep(base_params, k_values, thetaTD_values, y_o, v_o,
423                   results_path=RESULTS_PATH, rerun=RERUN_SWEEP):
424     if not rerun and os.path.exists(results_path):
425         data = np.load(results_path)
426         cached_k = data['k_values']
427         cached_theta = data['thetaTD_values']
428         if (cached_k.shape == k_values.shape and np.allclose(cached_k,
429                                                                k_values)
430             and cached_theta.shape == thetaTD_values.shape
431             and np.allclose(cached_theta, thetaTD_values)):
432             print(f"Loaded cached sweep from {results_path}")
433             return data['y_star_grid'], data['slope_grid']
434         else:
435             print("Cached grid does not match requested k_values/
436                   thetaTD_values - rerunning.")
437
438     print("Running fresh sweep...")
439     y_star_grid, slope_grid = gridSearchFixedPoints(base_params, k_values,
440                                                     thetaTD_values, y_o, v_o)
441     np.savez(results_path,
442              k_values=k_values,
443              thetaTD_values=thetaTD_values,
444              y_star_grid=y_star_grid,
445              slope_grid=slope_grid)
446     print(f"Saved sweep to {results_path}")
447     return y_star_grid, slope_grid
448
449 if __name__ == "__main__":
450     warnings.filterwarnings('error', category=np.exceptions.ComplexWarning)
451
452     # Bumped from 40x40: the local refinement check below showed the coarse
453     # grid's one-cell-wide stable "staircase" was a resolution artifact, not
454     # a real measure-zero band -- it thickened into a connected patch at 4x
455     # finer spacing. This resolution matches that refinement window's cell
456     # spacing (~0.26 deg, ~265 stiffness) across the full sweep range.
457     # ~12,600 cells; at ~0.47 s/cell (measured on this machine) that's
458     # roughly 90-100 minutes. Grab coffee.
459     k_sweep = np.linspace(5000, 40000, 133)
460     thetaTD_sweep_deg = np.linspace(55, 80, 95)
461     thetaTD_sweep_rad = np.radians(thetaTD_sweep_deg)

```

```

459
460 y_star_grid, slope_grid = runOrLoadSweep(params, k_sweep,
      thetaTD_sweep_rad, y_o, v_o)
461 print("y_star grid:\n", y_star_grid)
462 print("slope grid:\n", slope_grid)
463
464 fig, ax = plotStabilitySweep(k_sweep, thetaTD_sweep_deg, y_star_grid,
      slope_grid,
465                               friction_floor_deg=59.0,
466                               save_path=os.path.join(SCRIPT_DIR, '
      stability_sweep.png'))
467
468 coarse_row_run, coarse_col_run = stableBandWidth(slope_grid)
469 print(f"coarse grid stable band width: {coarse_row_run} cell(s) along
      theta, "
470       f"{coarse_col_run} cell(s) along k")
471
472 # ---- local refinement: zoom a 20x20 window (5 deg x 5000 stiffness,
473 # matching one dip below the friction-cone floor at 59 deg so we're not
474 # zooming into a physically-excluded corner) onto a stable cell from the
475 # coarse sweep and see if the band thickens at 4x the resolution ----
476 stable_coarse = (~np.isnan(slope_grid)) & (np.abs(slope_grid) < 1)
477 stable_idx = np.argwhere(stable_coarse)
478 if stable_idx.size == 0:
479     print("no stable cells found in the coarse sweep -- skipping
      refinement")
480 else:
481     i, j = stable_idx[len(stable_idx) // 2] # a representative stable
      cell, not an edge case
482     k_center = k_sweep[i]
483     theta_center_deg = thetaTD_sweep_deg[j]
484     print(f"refining around k={k_center:.1f}, thetaTD={theta_center_deg:.2
      f} deg")
485
486     k_fine, theta_fine_deg, y_star_fine, slope_fine = localRefinementSweep
      (
487         params, k_center, theta_center_deg,
488         k_half_width=2500.0, theta_half_width_deg=2.5, n=20, y_o=y_o, v_o=
      v_o)
489
490     fine_row_run, fine_col_run = stableBandWidth(slope_fine)
491     print(f"refined grid stable band width: {fine_row_run} cell(s) along
      theta, "
492           f"{fine_col_run} cell(s) along k (out of {len(theta_fine_deg)}x{
      len(k_fine)} cells, "
493           f"~{(theta_fine_deg[1]-theta_fine_deg[0]):.3f} deg / {(k_fine
      [1]-k_fine[0]):.1f} stiffness spacing)")
494
495     if fine_row_run <= 1 and fine_col_run <= 1:
496         print("band stays one-cell-wide at 4x finer resolution -- looks
      like a real, "
497               "genuinely narrow stability margin, not a resolution
      artifact.")
498     else:
499         print("band thickens at finer resolution -- the coarse sweep's
      staircase was ")

```

```
500         "a grid artifact; bump the full sweep's resolution before  
501           finalizing.")  
502     plotStabilitySweep(k_fine, theta_fine_deg, y_star_fine, slope_fine,  
503                       friction_floor_deg=59.0,  
504                       save_path=os.path.join(SCRIPT_DIR, '  
           stability_refinement.png'))
```